

UCS2612 Machine Learning Laboratory

Assignment1 : Working with Python packages - Numpy, Scipy, Scikit-learn, Matplotlib

Name : **Mega V**

Roll No : **3122 21 5001 051**

Aim :

To learning the python packages and working with that and to Explore the steps involved in the Learning process. For a dataset

Explore the various functions / methods that come under the following Python Libraries

a) **Numpy** – Numerical Python

Used for working with arrays.

Functions:

- Create a NumPy array

```
import numpy as np
arr = np.array([1, 2, 3, 4, 5])
print(arr)
print(type(arr))
```

```
[1 2 3 4 5]
<class 'numpy.ndarray'>
```

- Array Slicing

```
import numpy as np
arr = np.array([1, 2, 3, 4, 5, 6, 7])
print(arr[4:])
```

```
[5 6 7]
```

- Splitting is reverse operation of Joining

```
import numpy as np
arr = np.array([1, 2, 3, 4, 5, 6])
newarr = np.array_split(arr, 4)
print(newarr)
```

```
[array([1, 2]), array([3, 4]), array([5]), array([6])]
```

- Sorting arrays

```
import numpy as np
arr1 = np.array([3, 2, 0, 1])
arr2 = np.array(['banana', 'cherry', 'apple'])
arr3 = np.array([True, False, True])
print(np.sort(arr1))
print(np.sort(arr2))
print(np.sort(arr3))
```

```
[0 1 2 3]
['apple' 'banana' 'cherry']
[False  True  True]
```

- Filtering Arrays

```
import numpy as np
arr = np.array([41, 42, 43, 44])
x = [True, False, True, False]
newarr = arr[x]
print(newarr)
```

```
[41 43]
```

- Access Array Elements

```
import numpy as np
arr = np.array([[1,2,3,4,5], [6,7,8,9,10]])
print('2nd element on 1st row: ', arr[0, 1])
```

```
2nd element on 1st dim: 2
```

- Copy

```
import numpy as np
arr = np.array([1, 2, 3, 4, 5])
x = arr.copy()
arr[0] = 42
print(arr)
print(x)
```

```
[42 2 3 4 5]
[1 2 3 4 5]
```

- View

```
import numpy as np
arr = np.array([1, 2, 3, 4, 5])
x = arr.view()
arr[0] = 42
print(arr)
print(x)
```

```
[42 2 3 4 5]
[42 2 3 4 5]
```


- Make changes in view

```
import numpy as np
arr = np.array([1, 2, 3, 4, 5])
x = arr.view()
x[0] = 31
print(arr)
print(x)
```

```
[31 2 3 4 5]
[31 2 3 4 5]
```

- Seaborn - Seaborn is a library that uses Matplotlib underneath to plot graphs. It will be used to visualize random distributions.

```
import matplotlib.pyplot as plt
import seaborn as sns
sns.distplot([0, 1, 2, 3, 4, 5])
plt.show()
```



- Normal distribution
 - Use the `random.normal()` method to get a Normal Data Distribution.
 - It has three parameters:
 - `loc` - (Mean) where the peak of the bell exists.
 - `scale` - (Standard Deviation) how flat the graph distribution should be.
 - `size` - The shape of the returned array.


```
from numpy import random
x = random.normal(loc=1, scale=2, size=(2, 3))
print(x)
```

```
[[0.78604046 3.64592481 4.45391188]
 [0.68918731 1.22383744 4.65969041]]
```

```
from numpy import random
x = random.normal(size=(2, 3))
print(x)
```

```
[[ 0.70074132  0.24760778 -0.52154205]
 [-0.930514   -0.0909822  -0.95816753]]
```


```
from numpy import random
import matplotlib.pyplot as plt
import seaborn as sns
sns.distplot(random.normal(size=1000), hist=False)
plt.show()
```




- Binomial Distribution
 - It describes the outcome of binary scenarios, e.g. toss of a coin, it will either be head or tails.

- It has three parameters:
 - **n** - number of trials.
 - **p** - probability of occurrence of each trial (e.g. for toss of a coin 0.5 each).
 - **size** - The shape of the returned array.

```
from numpy import random
x = random.binomial(n=10, p=0.5, size=10)
print(x)
[7 5 7 9 6 4 6 4 3 5]
from numpy import random
import matplotlib.pyplot as plt
import seaborn as sns
sns.distplot(random.binomial(n=10, p=0.5, size=1000),
hist=True, kde=False)
plt.show()
```



```
from numpy import random
import matplotlib.pyplot as plt
import seaborn as sns
sns.distplot(random.normal(loc=50, scale=5,
size=1000), hist=False, label='normal')
sns.distplot(random.binomial(n=100, p=0.5,
size=1000), hist=False, label='binomial')
plt.show()
```



- Poisson Distribution


- It estimates how many times an event can happen in a specified time. e.g. If someone eats twice a day what is the probability he will eat thrice?
- It has two parameters:

- **lam** - rate or known number of occurrences e.g. 2 for above problem.
- **size** - The shape of the returned array.

```
from numpy import random
x = random.poisson(lam=2, size=10)
print(x)
```

```
[1 5 4 1 3 0 0 1 3 5]
```

```
from numpy import random
import matplotlib.pyplot as plt
import seaborn as sns
sns.distplot(random.poisson(lam=2, size=1000),
kde=False)
plt.show()
```




- Uniform Distribution

- Used to describe probability where every event has equal chances of occurring.
- E.g. Generation of random numbers.
- It has three parameters:
 - **a** - lower bound - default 0 .0.
 - **b** - upper bound - default 1.0.
 - **size** - The shape of the returned array.

```
from numpy import random
x = random.uniform(size=(2, 3))
print(x)
```

```
[[0.8300512  0.23174099 0.13371372]
 [0.56232815 0.84961613 0.58218586]]
```

```
from numpy import random
import matplotlib.pyplot as plt
import seaborn as sns
sns.distplot(random.uniform(size=1000), hist=False)
plt.show()
```



- Exponential Distribution
 - Exponential distribution is used for describing time till next event e.g. failure/success etc.
 - It has two parameters:
 - **scale** - inverse of rate (see lam in poisson distribution) defaults to 1.0.
 - **size** - The shape of the returned array.

```
from numpy import random
x = random.exponential(scale=2, size=(2, 3))
print(x)
```

```
[[2.08340211 3.71240206 0.21520718]
 [1.78793833 4.17123932 5.64152689]]
```

- Summations

```
import numpy as np
arr1 = np.array([1, 2, 3])
arr2 = np.array([1, 2, 3])
newarr = np.add(arr1, arr2)
print(newarr)
```

```
[2 4 6]
```

```
import numpy as np
arr1 = np.array([1, 2, 3])
arr2 = np.array([1, 2, 3])
newarr = np.sum([arr1, arr2])
print(newarr)
```

```
12
```

```
import numpy as np
arr = np.array([1, 2, 3])
newarr = np.cumsum(arr)
print(newarr)
```

```
[1 3 6]
```

- Trigonometric

```
import numpy as np
arr = np.array([np.pi/2, np.pi/3, np.pi/4, np.pi/5])
x = np.sin(arr)
print(x)
```

```
[1.          0.8660254  0.70710678 0.58778525]
```

b) Pandas

Used to analyze data.

Functions:

- Create an alias with the `as` keyword

```
import pandas as pd
mydataset = {
    'cars': ["BMW", "Volvo", "Ford"],
    'passings': [3, 7, 2]
}
myvar = pd.DataFrame(mydataset)
print(myvar)
```

	cars	passings
0	BMW	3
1	Volvo	7
2	Ford	2

- Create simple pandas series from the list.

```
import pandas as pd
a = [1, 7, 2]
myvar = pd.Series(a)
print(myvar)
```

0	1
1	7
2	2

dtype: int64

- Index argument

```
import pandas as pd
a = [1, 7, 2]
myvar = pd.Series(a, index = ["x", "y", "z"])
print(myvar)
```

x	1
y	7
z	2

dtype: int64

```
import pandas as pd
calories = {"day1": 420, "day2": 380, "day3": 390}
myvar = pd.Series(calories, index = ["day1", "day2"])
print(myvar)
```

day1	420
day2	380

dtype: int64

- DataFrame from two series

```
import pandas as pd
data = {
    "calories": [420, 380, 390],
    "duration": [50, 40, 45]
}
df = pd.DataFrame(data)
print(myvar)
```

	calories	duration
0	420	50
1	380	40
2	390	45

- Locate row

```
print(df.loc[0])
```

calories	420
duration	50
Name: 0, dtype: int64	

```
print(df.loc[[0, 1]])
```

	calories	duration
0	420	50
1	380	40

- Named indexes

```
import pandas as pd
data = {
    "calories": [420, 380, 390],
    "duration": [50, 40, 45]
}
df = pd.DataFrame(data, index =
["day1", "day2", "day3"])
print(df)
```

	calories	duration
day1	420	50
day2	380	40
day3	390	45

- Load files into a dataframe

```
import pandas as pd
df = pd.read_csv('data.csv')
print(df)
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0
..
164	60	105	140	290.8
165	60	110	145	300.4
166	60	115	145	310.2
167	75	120	150	320.4
168	75	125	150	330.4

[169 rows x 4 columns]

- Max rows

```
import pandas as pd
print(pd.options.display.max_rows)
```

```
60
```

```
import pandas as pd
pd.options.display.max_rows = 6
df = pd.read_csv('data.csv')
print(df)
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
...
166	60	115	145	310.2
167	75	120	150	320.4
168	75	125	150	330.4

[169 rows x 4 columns]

- Read JSON

```
import pandas as pd
df = pd.read_json('data.json')
print(df.to_string())
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0
5	60	102	127	300.5
6	60	110	136	374.0
7	45	104	134	253.3
8	30	109	133	195.1
9	60	98	124	269.0
10	60	103	147	329.3

- Dictionary as JSON

```
import pandas as pd
data = {
    "Duration":{
        "0":60,
        "1":60,
        "2":60,
        "3":45,
        "4":45,
        "5":60
    },
    "Pulse":{
        "0":110,
        "1":117,
        "2":103,
        "3":109,
```

```

        "4":117,
        "5":102
    },
    "Maxpulse":{
        "0":130,
        "1":145,
        "2":135,
        "3":175,
        "4":148,
        "5":127
    },
    "Calories":{
        "0":409,
        "1":479,
        "2":340,
        "3":282,
        "4":406,
        "5":300
    }
}
df = pd.DataFrame(data)
print(df)

```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0
5	60	102	127	300.5

- Viewing data – first 5 rows, last 5 rows

```

import pandas as pd
df = pd.read_csv('data.csv')
print(df.head())
print(df.tail())
print(df.info())

```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0

	Duration	Pulse	Maxpulse	Calories
164	60	105	140	290.8
165	60	110	145	300.4
166	60	115	145	310.2
167	75	120	150	320.4
168	75	125	150	330.4

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 169 entries, 0 to 168
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Duration    169 non-null    int64
1   Pulse       169 non-null    int64
2   Maxpulse    169 non-null    int64
3   Calories    164 non-null    float64
dtypes: float64(1), int64(3)
memory usage: 5.4 KB
None
```

- Data cleaning
 - o Data cleaning means fixing bad data in your data set. Bad data could be:
 - Empty cells
 - Data in wrong format
 - Wrong data
 - Duplicates

Cleaning empty cells (has NaN values) - Remove rows with NaN values:

```
import pandas as pd
df = pd.read_csv('data.csv')
new_df = df.dropna()
print(new_df.to_string())
```

Data of wrong format converted into correct format:

19	60	'2020/12/19'	103	123	323.0
20	45	'2020/12/20'	97	125	243.0
21	60	'2020/12/21'	108	131	364.2
22	45	NaN	100	119	282.0
23	60	'2020/12/23'	130	101	300.0
24	45	'2020/12/24'	105	132	246.0
25	60	'2020/12/25'	102	126	334.5
26	60	20201226	100	120	250.0
27	60	'2020/12/27'	92	118	241.0

```
import pandas as pd
df = pd.read_csv('data.csv')
df['Date'] = pd.to_datetime(df['Date'])
print(df.to_string())
```

19	60	'2020/12/19'	103	123	323.0
20	45	'2020/12/20'	97	125	243.0
21	60	'2020/12/21'	108	131	364.2
22	45	NaT	100	119	282.0
23	60	'2020/12/23'	130	101	300.0
24	45	'2020/12/24'	105	132	246.0
25	60	'2020/12/25'	102	126	334.5
26	60	'2020/12/26'	100	120	250.0
27	60	'2020/12/27'	92	118	241.0

```
df.dropna(subset=['Date'], inplace = True)
```

19	60	2020-12-19	103	123	323.0
20	45	2020-12-20	97	125	243.0
21	60	2020-12-21	108	131	364.2
23	60	2020-12-23	130	101	300.0
24	45	2020-12-24	105	132	246.0
25	60	2020-12-25	102	126	334.5
26	60	2020-12-26	100	120	250.0
27	60	2020-12-27	92	118	241.0

- Fixing wrong data

5	60	'2020/12/06'	102	127	300.0
6	60	'2020/12/07'	110	136	374.0
7	450	'2020/12/08'	104	134	253.3
8	30	'2020/12/09'	109	133	195.1
9	60	'2020/12/10'	98	124	269.0
10	60	'2020/12/11'	103	147	329.3

In our example, it is most likely a typo, and the value should be "45" instead of "450", and we could just insert "45" in row 7:

```
import pandas as pd
df = pd.read_csv('data.csv')
df.loc[7, 'Duration'] = 45
print(df.to_string())
```

5	60	'2020/12/06'	102	127	300.0
6	60	'2020/12/07'	110	136	374.0
7	45	'2020/12/08'	104	134	253.3
8	30	'2020/12/09'	109	133	195.1
9	60	'2020/12/10'	98	124	269.0
10	60	'2020/12/11'	103	147	329.3

- Discover duplicates

9	60	'2020/12/10'	98	124	269.0
10	60	'2020/12/11'	103	147	329.3
11	60	'2020/12/12'	100	120	250.7
12	60	'2020/12/12'	100	120	250.7
13	60	'2020/12/13'	106	128	345.3

```
print(df.duplicated())
```

9	False
10	False
11	False
12	True
13	False


- Remove duplicates
df.drop_duplicates(inplace = True)

9	60	'2020/12/10'	98	124	269.0
10	60	'2020/12/11'	103	147	329.3
11	60	'2020/12/12'	100	120	250.7
13	60	'2020/12/13'	106	128	345.3

- Finding relationships
df.corr()

	Duration	Pulse	Maxpulse	Calories
Duration	1.000000	-0.155408	0.009403	0.922721
Pulse	-0.155408	1.000000	0.786535	0.025120
Maxpulse	0.009403	0.786535	1.000000	0.203814
Calories	0.922721	0.025120	0.203814	1.000000


- Plotting
import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv('data.csv')
df.plot()
plt.show()



- Scatter plot

```
import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv('data.csv')
df.plot(kind = 'scatter', x = 'Duration',
              y = 'Calories')

plt.show()
```



c) SciPy

Working with graphs

Graphs are an essential data structure.

SciPy provides us with the module `scipy.sparse.csgraph` for working with such data structures.

SciPy is built on NumPy and provides additional functionality for scientific computing.

It includes modules for optimization, integration, interpolation, eigenvalue problems, and more. Adjacency matrix


- Key Functions/Methods:

- Integration: `scipy.integrate.quad()`, `scipy.integrate.simps()`.
- Interpolation: `scipy.interpolate.interpld()`.
- Optimization: `scipy.optimize.minimize()`.
- Signal Processing: `scipy.signal.convolve()`,
`scipy.signal.spectrogram()`.
- Special functions: `scipy.special.erf()`, `scipy.special.gamma()`.

```

      A B C
A: [0 1 2]
B: [1 0 0]
C: [2 0 0]

```



- Connected components


```

import numpy as np
from scipy.sparse.csgraph import connected_components
from scipy.sparse import csr_matrix
arr = np.array([
    [0, 1, 2],
    [1, 0, 0],
    [2, 0, 0]
])
newarr = csr_matrix(arr)
print(connected_components(newarr))
(1, array([0, 0, 0], dtype=int32))

```
- Dijkstra


```

import numpy as np
from scipy.sparse.csgraph import dijkstra
from scipy.sparse import csr_matrix
arr = np.array([
    [0, 1, 2],
    [1, 0, 0],
    [2, 0, 0]
])
newarr = csr_matrix(arr)
print(dijkstra(newarr, return_predecessors=True,
indices=0))
(array([ 0., 1., 2.]), array([-9999, 0, 0], dtype=int32))

```
- Floyd warshall


```

import numpy as np
from scipy.sparse.csgraph import floyd_warshall
from scipy.sparse import csr_matrix
arr = np.array([
    [0, 1, 2],
    [1, 0, 0],
    [2, 0, 0]
])
newarr = csr_matrix(arr)

```

```
print(floyd_warshall(newarr,
return_predecessors=True))
(array([[ 0.,  1.,  2.],
       [ 1.,  0.,  3.],
       [ 2.,  3.,  0.])), array([[ -9999,    0,    0],
       [    1, -9999,    0],
       [    2,    0, -9999]], dtype=int32))
```

- Bellman Ford

```
import numpy as np
from scipy.sparse.csgraph import bellman_ford
from scipy.sparse import csr_matrix
arr = np.array([
    [0, -1, 2],
    [1, 0, 0],
    [2, 0, 0]
])
newarr = csr_matrix(arr)
print(bellman_ford(newarr, return_predecessors=True,
indices=0))
(array([ 0., -1.,  2.]), array([ -9999,    0,    0], dtype=int32))
```

- Depth first order

```
import numpy as np
from scipy.sparse.csgraph import depth_first_order
from scipy.sparse import csr_matrix
arr = np.array([
    [0, 1, 0, 1],
    [1, 1, 1, 1],
    [2, 1, 1, 0],
    [0, 1, 0, 1]
])
newarr = csr_matrix(arr)
print(depth_first_order(newarr, 1))
(array([1, 0, 3, 2], dtype=int32), array([ 1, -9999,  1,  0], dtype=int32))
```

- Breadth first order

```
import numpy as np
from scipy.sparse.csgraph import breadth_first_order
from scipy.sparse import csr_matrix
arr = np.array([
    [0, 1, 0, 1],
    [1, 1, 1, 1],
    [2, 1, 1, 0],
    [0, 1, 0, 1]
])
newarr = csr_matrix(arr)
print(breadth_first_order(newarr, 1))
(array([1, 0, 2, 3], dtype=int32), array([ 1, -9999,  1,  1], dtype=int32))
```

- Binary Entropy Function

```
def binary_entropy(x):
    return -(sc.xlogy(x, x) + sc.xlog1py(1 - x, -
x))/np.log(2)
```


- A rectangular step function on [0,1]

```
def step(x):
    return 0.5*(np.sign(x) + np.sign(1 - x))
```

- Ramp function


```
def ramp(x):
    return np.maximum(0, x)
```

- Bessel functions of real order (jv, jn_zeros)

Bessel functions are a family of solutions to Bessel's differential equation with real or complex order alpha:

$$x^2 \frac{d^2 y}{dx^2} + x \frac{dy}{dx} + (x^2 - \alpha^2)y = 0$$

```
>>> from scipy import special
>>> import numpy as np
>>> def drumhead_height(n, k, distance, angle, t):
...     kth_zero = special.jn_zeros(n, k)[-1]
...     return np.cos(t) * np.cos(n*angle) *
special.jn(n, distance*kth_zero)
>>> theta = np.r_[0:2*np.pi:50j]
>>> radius = np.r_[0:1:50j]
>>> x = np.array([r * np.cos(theta) for r in
radius])
>>> y = np.array([r * np.sin(theta) for r in
radius])
>>> z = np.array([drumhead_height(1, 1, r, theta,
0.5) for r in radius])
>>> import matplotlib.pyplot as plt
>>> fig = plt.figure()
>>> ax = fig.add_axes(rect=(0, 0.05, 0.95, 0.95),
projection='3d')
>>> ax.plot_surface(x, y, z, rstride=1, cstride=1,
cmap='RdBu_r', vmin=-0.5, vmax=0.5)
>>> ax.set_xlabel('X')
>>> ax.set_ylabel('Y')
>>> ax.set_xticks(np.arange(-1, 1.1, 0.5))
>>> ax.set_yticks(np.arange(-1, 1.1, 0.5))
>>> ax.set_zlabel('Z')
>>> plt.show()
```



- Cython Bindings for Special Functions (scipy.special.cython_special)

```
cimport scipy.special.cython_special as csc
cdef:
    double x = 1
    double complex z = 1 + 1j
    double si, ci, rgam
    double complex cgam
rgam = csc.gamma(x)
print(rgam)
cgam = csc.gamma(z)
print(cgam)
csc.sici(x, &si, &ci)
print(si, ci)
```

- Nelder-Mead Simplex algorithm(method= 'Nelder-Mead')

```
>>> import numpy as np
>>> from scipy.optimize import minimize
```

```
>>> def rosen(x):
...     """The Rosenbrock function"""
...     return sum(100.0*(x[1:]-x[:-1]**2.0)**2.0 + (1-x[:-1])**2.0)
```

```
>>> x0 = np.array([1.3, 0.7, 0.8, 1.9, 1.2])
>>> res = minimize(rosen, x0, method='nelder-mead',
...                 options={'xatol': 1e-8, 'disp': True})
Optimization terminated successfully.
Current function value: 0.000000
Iterations: 339
Function evaluations: 571
```

```
>>> print(res.x)
[1. 1. 1. 1. 1.]
```

d) Scikit-Learn

Identify which category an object belongs to.


Applications: spam detection, image recognition

Algorithms: [Gradient boosting](#), [nearest neighbors](#), [random forest](#), [logistic regression](#), classification, regression, clustering, and dimensionality reduction.

- Regression

Applications: drug response, stock prices.

Algorithms: [Gradient boosting](#), [nearest neighbors](#), [random forest](#), [ridge](#), and [more...](#)




- Clustering

Applications : Customer segmentation, Grouping experiment outcomes

Algorithms: [k-Means](#), [HDBSCAN](#), [hierarchical clustering](#), and [more...](#)

K-means clustering on the digits dataset (PCA-reduced data)
Centroids are marked with white cross




- Dimensionality Reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: [PCA](#), [feature selection](#), [non-negative matrix factorization](#), and [...](#)




- Model selection

Comparing, validating and choosing parameters and models.

Applications: Improved accuracy via parameter tuning.

Algorithms: [grid search](#), [cross validation](#), [metrics](#), and [more...](#)



- Preprocessing

Feature extraction and normalization.

Applications: Transforming input data such as text for use with machine learning algorithms.

Algorithms: [preprocessing](#), [feature extraction](#), and [more...](#)

- Key Functions/Methods:

- `sklearn.model_selection.train_test_split()`: Split data into training and testing sets.
- Model building and training: `fit()`, `predict()`.
- Evaluation metrics: `sklearn.metrics.accuracy_score()`, `sklearn.metrics.confusion_matrix()`.
- Preprocessing: `sklearn.preprocessing.StandardScaler()`, `sklearn.preprocessing.LabelEncoder()`.

- Fitting and predicting: estimator basics

- Scikit-learn provides dozens of built-in machine learning algorithms and models, called estimators.
- Each estimator can be fitted to some data using its fit method.
 - `RandomForestClassifier`

```
from sklearn.ensemble import
RandomForestClassifier
clf = RandomForestClassifier(random_state=0)
X = [[ 1,  2,  3], # 2 samples, 3 features
```

```
...      [11, 12, 13]]
y = [0, 1] # classes of each sample
clf.fit(X, y)
RandomForestClassifier(random_state=0)
```

Once the estimator is fitted, it can be used for predicting target values of new data. You don't need to re-train the estimator:

```
clf.predict(X) # predict classes of the
               training data
array([0, 1])
clf.predict([[4, 5, 6], [14, 15, 16]]) #
               predict classes of new data
array([0, 1])
```

- Transformers and pre-processors

Pre-processors and transformers follow the same API as the estimator objects (they actually all inherit from the same `BaseEstimator` class). The transformer objects don't have a `predict` method but rather a `transform` method that outputs a newly transformed sample matrix `X`:

```
>>> from sklearn.preprocessing import
StandardScaler
>>> X = [[0, 15],
...      [1, -10]]
>>> # scale data according to computed scaling
values
>>> StandardScaler().fit(X).transform(X)
array([[ -1.,   1.],
       [  1.,  -1.]])
```

- Pipelines: chaining pre-processors and estimators

Transformers and estimators (predictors) can be combined together into a single unifying object: a Pipeline.

```
>>> from sklearn.preprocessing import StandardScaler
>>> from sklearn.linear_model import LogisticRegression
>>> from sklearn.pipeline import make_pipeline
>>> from sklearn.datasets import load_iris
>>> from sklearn.model_selection import train_test_split
>>> from sklearn.metrics import accuracy_score
...
>>> # create a pipeline object
>>> pipe = make_pipeline(
...     StandardScaler(),
...     LogisticRegression()
... )
...
>>> # load the iris dataset and split it into train and
test sets
>>> X, y = load_iris(return_X_y=True)
>>> X_train, X_test, y_train, y_test = train_test_split(X,
y, random_state=0)
...

```



```

>>> # now create a searchCV object and fit it to the data
>>> search =
RandomizedSearchCV(estimator=RandomForestRegressor(random_state=0),
...                 n_iter=5,
...
param_distributions=param_distributions,
...                 random_state=0)
>>> search.fit(X_train, y_train)
RandomizedSearchCV(estimator=RandomForestRegressor(random_state=0), n_iter=5,
...                 param_distributions={'max_depth': ...,
...                                     'n_estimators': ...},
...                 random_state=0)
>>> search.best_params_
{'max_depth': 9, 'n_estimators': 4}

>>> # the search object now acts like a normal random forest
estimator
>>> # with max_depth=9 and n_estimators=4
>>> search.score(X_test, y_test)
0.73...

```

e) Matplotlib

Matplotlib is a 2D plotting library for creating static, animated, and interactive visualizations in Python.

- Key Functions/Methods:
 - matplotlib.pyplot.plot(), matplotlib.pyplot.scatter(): Create line plots and scatter plots.
 - matplotlib.pyplot.xlabel(), matplotlib.pyplot.ylabel(): Set axis labels.
 - matplotlib.pyplot.legend(): Add legends to plots.
 - matplotlib.pyplot.subplot(), matplotlib.pyplot.figure(): Create subplots and figures.
 - matplotlib.pyplot.savefig(): Save figures to a file.

```


import matplotlib.pyplot as plt
import numpy as np

xpoints = np.array([0, 6])
ypoints = np.array([0, 250])

plt.plot(xpoints, ypoints)

```

```
plt.show()
```




- Plotting x and y points

```
import matplotlib.pyplot as plt  
import numpy as np
```


```
xpoints = np.array([1, 8])  
ypoints = np.array([3, 10])
```

```
plt.plot(xpoints, ypoints)  
plt.show()
```




- Plotting Without Line

```
import matplotlib.pyplot as plt  
import numpy as np  
xpoints = np.array([1, 8])  
ypoints = np.array([3, 10])  
plt.plot(xpoints, ypoints, 'o')  
plt.show()
```


- Multiple Points

```
import matplotlib.pyplot as plt
import numpy as np
xpoints = np.array([1, 2, 6, 8])
ypoints = np.array([3, 8, 1, 10])
plt.plot(xpoints, ypoints)
plt.show()
```




- Markers

```
import matplotlib.pyplot as plt
import numpy as np
ypoints = np.array([3, 8, 1, 10])
plt.plot(ypoints, marker = 'o')
plt.show()
```




```
plt.plot(ypoints, marker = '*')
```




- Format Strings fmt

```
import matplotlib.pyplot as plt
import numpy as np
ypoints = np.array([3, 8, 1, 10])
plt.plot(ypoints, 'o:r')
plt.show()
```




- Marker Size

```
import matplotlib.pyplot as plt
import numpy as np
ypoints = np.array([3, 8, 1, 10])
plt.plot(ypoints, marker = 'o', ms = 20)
plt.show()
```




- Marker Color


```
import matplotlib.pyplot as plt
import numpy as np
ypoints = np.array([3, 8, 1, 10])
plt.plot(ypoints, marker = 'o', ms = 20, mec = 'r')
plt.show()
```



```
import matplotlib.pyplot as plt
import numpy as np
ypoints = np.array([3, 8, 1, 10])
plt.plot(ypoints, marker = 'o', ms = 20, mfc = 'r')
plt.show()
```




```
import matplotlib.pyplot as plt
import numpy as np
ypoints = np.array([3, 8, 1, 10])
plt.plot(ypoints, marker = 'o', ms = 20, mec = 'r',
mfc = 'r')
plt.show()
```




- Linestyle

```
import matplotlib.pyplot as plt
import numpy as np
ypoints = np.array([3, 8, 1, 10])
plt.plot(ypoints, linestyle = 'dotted') # "dashed"
plt.show()
```




- Shorter syntax

```
plt.plot(ypoints, ls = ':')
```



- Line color

```
import matplotlib.pyplot as plt
import numpy as np
ypoints = np.array([3, 8, 1, 10])
plt.plot(ypoints, color = 'r') # c='#4CAF50', c='hotpink'
plt.show()
```




- Line width

```
import matplotlib.pyplot as plt
import numpy as np


ypoints = np.array([3, 8, 1, 10])

plt.plot(ypoints, linewidth = '20.5')
plt.show()
```



- Multiple Lines


```
import matplotlib.pyplot as plt
import numpy as np
y1 = np.array([3, 8, 1, 10])
y2 = np.array([6, 2, 7, 11])
plt.plot(y1)
plt.plot(y2)
plt.show()
```



- Set font properties for title and labels


```
import numpy as np
import matplotlib.pyplot as plt
x =
np.array([80, 85, 90, 95, 100, 105, 110, 115, 120, 12
5])
y =
np.array([240, 250, 260, 270, 280, 290, 300, 310, 320
, 330])
```

```
font1 = {'family':'serif','color':'blue','size':20}
font2 =
{'family':'serif','color':'darkred','size':15}
plt.title("Sports Watch Data", fontdict = font1)
plt.xlabel("Average Pulse", fontdict = font2)
plt.ylabel("Calorie Burnage", fontdict = font2)
plt.plot(x, y)
plt.show()
```



- Position of the title

```
import numpy as np
import matplotlib.pyplot as plt
x =
np.array([80, 85, 90, 95, 100, 105, 110, 115, 120, 125])
y =
np.array([240, 250, 260, 270, 280, 290, 300, 310, 320, 330])
plt.title("Sports Watch Data", loc = 'left')
plt.xlabel("Average Pulse")
plt.ylabel("Calorie Burnage")
plt.plot(x, y)
plt.show()
```



- Grid lines

```
import numpy as np
import matplotlib.pyplot as plt


x =
np.array([80, 85, 90, 95, 100, 105, 110, 115, 120, 125])
y =
np.array([240, 250, 260, 270, 280, 290, 300, 310, 320, 330])

plt.title("Sports Watch Data")
plt.xlabel("Average Pulse")
plt.ylabel("Calorie Burnage")

plt.plot(x, y)

plt.grid(color = 'green', linestyle = '--', linewidth = 0.5)

plt.show()
```



- Subplot for 6 charts

```
import matplotlib.pyplot as plt
import numpy as np

x = np.array([0, 1, 2, 3])
y = np.array([3, 8, 1, 10])

plt.subplot(2, 3, 1)
plt.plot(x,y)

x = np.array([0, 1, 2, 3])
y = np.array([10, 20, 30, 40])

plt.subplot(2, 3, 2)
plt.plot(x,y)
```

```

x = np.array([0, 1, 2, 3])
y = np.array([3, 8, 1, 10])

plt.subplot(2, 3, 3)
plt.plot(x,y)

x = np.array([0, 1, 2, 3])
y = np.array([10, 20, 30, 40])

plt.subplot(2, 3, 4)
plt.plot(x,y)

x = np.array([0, 1, 2, 3])
y = np.array([3, 8, 1, 10])


plt.subplot(2, 3, 5)
plt.plot(x,y)

x = np.array([0, 1, 2, 3])
y = np.array([10, 20, 30, 40])

plt.subplot(2, 3, 6)
plt.plot(x,y)

plt.show()

```



- Subtitle

```

import matplotlib.pyplot as plt
import numpy as np

#plot 1:
x = np.array([0, 1, 2, 3])
y = np.array([3, 8, 1, 10])

plt.subplot(1, 2, 1)
plt.plot(x,y)
plt.title("SALES")

```



```
#plot 2:
x = np.array([0, 1, 2, 3])
y = np.array([10, 20, 30, 40])

plt.subplot(1, 2, 2)
plt.plot(x,y)
plt.title("INCOME")

plt.suptitle("MY SHOP")
plt.show()
```




- Scatter Plots()

```
import matplotlib.pyplot as plt
import numpy as np

x = np.array([5,7,8,7,2,17,2,9,4,11,12,9,6])
y =
np.array([99,86,87,88,111,86,103,87,94,78,77,85,86])

plt.scatter(x, y)
plt.show()
```




- Compare plots

```
import matplotlib.pyplot as plt
import numpy as np
```

A scatter plot showing the relationship between the number of children (X-axis) and the number of books read (Y-axis). The X-axis ranges from 0 to 16, and the Y-axis ranges from 75 to 115. Data points are colored blue and orange, showing a general downward trend.

Number of Children (X)	Number of Books Read (Y)	Color
1	105	Orange
2	111	Blue
2	105	Orange
2	100	Orange
3	100	Orange
4	112	Orange
4	94	Blue
5	99	Blue
6	86	Blue
7	86	Blue
7	88	Blue
7	91	Orange
7	94	Orange
8	84	Orange
8	87	Blue
8	99	Orange
9	87	Blue
9	95	Orange
11	78	Blue
11	79	Orange
12	77	Blue
12	85	Orange
12	90	Orange
14	80	Orange
15	90	Orange
17	86	Blue


- ```
import matplotlib.pyplot as plt
import numpy as np
x = np.array([5,7,8,7,2,17,2,9,4,11,12,9,6])
y =
np.array([99,86,87,88,111,86,103,87,94,78,77,85,86])
sizes
= np.array([20,50,100,200,500,1000,60,90,10,300,600,8
00,75])
plt.scatter(x, y, s=sizes)
plt.show()
```



- Bars – For x-axis: bar(x,y), For y-axis: barh(x,y)  
`import matplotlib.pyplot as plt`  
`import numpy as np`

```
x = np.array(["A", "B", "C", "D"])
y = np.array([3, 8, 1, 10])
```


```
plt.bar(x,y)
plt.show()
```



- Bar color  
`import matplotlib.pyplot as plt`  
`import numpy as np`

```
x = np.array(["A", "B", "C", "D"])
y = np.array([3, 8, 1, 10])
```


```
plt.bar(x, y, color = "red")
plt.show()
```



- Histogram  
`import matplotlib.pyplot as plt`  
`import numpy as np`

```
x = np.random.normal(170, 10, 250)
```

```
plt.hist(x)
plt.show()
```




- Pie Chart

```
import matplotlib.pyplot as plt
import numpy as np
```

```
y = np.array([35, 25, 25, 15])
mylabels = ["Apples", "Bananas", "Cherries", "Dates"]
```

```
plt.pie(y, labels = mylabels)
plt.show()
```



List down the features and class labels from the dataset. Explore the steps involved in the Learning process.

- 1) Loading the dataset.
- 2) Pre-Processing the data (Handling missing values, Normalization, Standardization).
- 3) Exploratory Data Analysis.
- 4) Feature Selection Techniques.
- 5) Split the data into training, testing and validation sets.

**CODE and Output:**

```
import numpy as np
import pandas as pd
import scipy as sl
import matplotlib.pyplot as mat
```

Loading The Dataset

```
df = pd.read_csv("C:\\Users\\SSN\\Desktop\\ML
Lab\\Ex1\\diabetes_dataset.csv",)
```

```
print("The Shape of the Data Frame (rows , columns) : ",df.shape)
print()
print()
print("The DataType for Each Attributes in Data Frame\\n",df.dtypes)
```

```
The Shape of the Data Frame (rows , columns) : (100000, 9)
```

```
The DataType for Each Attributes in Data Frame
```

|                     |         |
|---------------------|---------|
| gender              | object  |
| age                 | float64 |
| hypertension        | float64 |
| heart_disease       | float64 |
| smoking_history     | object  |
| bmi                 | float64 |
| HbA1c_level         | float64 |
| blood_glucose_level | float64 |
| diabetes            | float64 |
| dtype:              | object  |

```
print(df.head)
```

```

... <bound method NDFrame.head of gender age hypertension heart_disease smoking_history bmi \
0 Female 80.0 0.0 1.0 never 25.19
1 Female 54.0 0.0 0.0 No Info 27.32
2 Male 28.0 0.0 0.0 never 27.32
3 Female 36.0 0.0 0.0 current 23.45
4 Male 76.0 1.0 1.0 current 20.14
...
99995 Female 80.0 0.0 0.0 No Info 27.32
99996 Female 2.0 0.0 0.0 No Info 17.37
99997 Male 66.0 0.0 0.0 former 27.83
99998 Female 24.0 0.0 0.0 never 35.42
99999 Female 57.0 0.0 0.0 current 22.43

 HbA1c_level blood_glucose_level diabetes
0 6.6 140.0 0.0
1 6.6 80.0 0.0
2 5.7 158.0 0.0
3 5.0 155.0 0.0
4 4.8 155.0 0.0
...
99995 6.2 90.0 0.0
99996 6.5 100.0 0.0
99997 5.7 155.0 0.0
99998 4.0 100.0 0.0
99999 6.6 90.0 0.0

[100000 rows x 9 columns]>

```

## Finding The Missing Values in the Dataframe

```
print("The Number of Missing Values in Each Columns\n",df.isnull().sum())
```

```

... The Number of Missing Values in Each Columns
gender 1
age 2
hypertension 2
heart_disease 1
smoking_history 16
bmi 14
HbA1c_level 7
blood_glucose_level 16
diabetes 4
dtype: int64

```

## Handling The Missing Values for The Attributes as follow

```

Gender Fill as Unknown
Age Fill a Median Value
Hypertension Fill a Null Value 0
Heart_disease Fill a Null value 0
Smoking History Fill a No Info

```

|                     |                     |
|---------------------|---------------------|
| BMI                 | Fill a Median Value |
| HbA1c_level         | Fill a Median Value |
| blood_glucose_level | Fill a Median Value |
| diabetes            | Remopve the Row     |

```
print("Missing values handling for 'gender'")
df['gender']=df['gender'].fillna('Unknown')
```

```
print("Missing values handling for Age")
age_median = df['age'].median()
df['age']=df['age'].fillna(age_median)
```

```
print("Missing values handling for Hyper Tension and Heart Disease")
df['hypertension']=df['hypertension'].fillna(0)
df['heart_disease']=df['heart_disease'].fillna(0)
```

```
print("Missing values handling for Smoking history")
df['smoking_history']=df['smoking_history'].fillna('No Info')
```

```
print("Missing values handling for bmi")
bmi_median = df['bmi'].median()
df['bmi']=df['bmi'].fillna(bmi_median)
```

```
print("Missing values handling for HbA1c_level")
hba1c_median = df['HbA1c_level'].median()
df['HbA1c_level']=df['HbA1c_level'].fillna(hba1c_median)
```

```
print("Missing values handling for HbA1c_level")
blood_glucose_median = df['blood_glucose_level'].median()
df['blood_glucose_level'] =
df['blood_glucose_level'].fillna(blood_glucose_median)
```

```
print("Missing values handling for diabetes_status")
df.dropna(subset=['diabetes'], inplace=True)
```

```
print("The Number of Missing Values in Each Columns\n",df.isnull().sum())
```

The Number of Missing Values in Each Columns

```
gender 0
age 0
hypertension 0
heart_disease 0
smoking_history 0
bmi 0
HbA1c_level 0
blood_glucose_level 0
diabetes 0
dtype: int64
```

Normalization The attributes Age and BMI

```
from sklearn.preprocessing import MinMaxScaler, StandardScaler
attributes_to_normalize = ['age', 'bmi']
min_max_scaler = MinMaxScaler()
df[attributes_to_normalize] =
min_max_scaler.fit_transform(df[attributes_to_normalize])
```

|       | gender | age      | hypertension | heart_disease | smoking_history | \ |
|-------|--------|----------|--------------|---------------|-----------------|---|
| 0     | Female | 1.000000 | 0.0          | 1.0           | never           |   |
| 1     | Female | 0.674675 | 0.0          | 0.0           | No Info         |   |
| 2     | Male   | 0.349349 | 0.0          | 0.0           | never           |   |
| 3     | Female | 0.449449 | 0.0          | 0.0           | current         |   |
| 4     | Male   | 0.949950 | 1.0          | 1.0           | current         |   |
| ...   | ...    | ...      | ...          | ...           | ...             |   |
| 99995 | Female | 1.000000 | 0.0          | 0.0           | No Info         |   |
| 99996 | Female | 0.024024 | 0.0          | 0.0           | No Info         |   |
| 99997 | Male   | 0.824825 | 0.0          | 0.0           | former          |   |
| 99998 | Female | 0.299299 | 0.0          | 0.0           | never           |   |
| 99999 | Female | 0.712212 | 0.0          | 0.0           | current         |   |

|       | bmi      | HbA1c_level | blood_glucose_level | diabetes |
|-------|----------|-------------|---------------------|----------|
| 0     | 0.263246 | 6.6         | 140.0               | 0.0      |
| 1     | 0.285505 | 6.6         | 80.0                | 0.0      |
| 2     | 0.285505 | 5.7         | 158.0               | 0.0      |
| 3     | 0.245062 | 5.0         | 155.0               | 0.0      |
| 4     | 0.210471 | 4.8         | 155.0               | 0.0      |
| ...   | ...      | ...         | ...                 | ...      |
| 99995 | 0.285505 | 6.2         | 90.0                | 0.0      |
| 99996 | 0.181524 | 6.5         | 100.0               | 0.0      |
| 99997 | 0.290835 | 5.7         | 155.0               | 0.0      |
| 99998 | 0.370154 | 4.0         | 100.0               | 0.0      |
| 99999 | 0.234403 | 6.6         | 90.0                | 0.0      |

[99996 rows x 9 columns]



## Standardization the Attributes HbA1c Level and Blood Glucose Level

```
List of attributes to be standardized
attributes_to_standardize = ['HbA1c_level', 'blood_glucose_level']

standard_scaler = StandardScaler()

df[attributes_to_standardize] =
standard_scaler.fit_transform(df[attributes_to_standardize])
```

|       | gender | age      | hypertension | heart_disease | smoking_history | \    |
|-------|--------|----------|--------------|---------------|-----------------|------|
| 0     | Female | 1.000000 | 0.0          | 1.0           | never           |      |
| 1     | Female | 0.674675 | 0.0          | 0.0           | No              | Info |
| 2     | Male   | 0.349349 | 0.0          | 0.0           | never           |      |
| 3     | Female | 0.449449 | 0.0          | 0.0           | current         |      |
| 4     | Male   | 0.949950 | 1.0          | 1.0           | current         |      |
| ...   | ...    | ...      | ...          | ...           | ...             | ...  |
| 99995 | Female | 1.000000 | 0.0          | 0.0           | No              | Info |
| 99996 | Female | 0.024024 | 0.0          | 0.0           | No              | Info |
| 99997 | Male   | 0.824825 | 0.0          | 0.0           | former          |      |
| 99998 | Female | 0.299299 | 0.0          | 0.0           | never           |      |
| 99999 | Female | 0.712212 | 0.0          | 0.0           | current         |      |

|       | bmi      | HbA1c_level | blood_glucose_level | diabetes |
|-------|----------|-------------|---------------------|----------|
| 0     | 0.263246 | 1.001694    | 0.047797            | 0.0      |
| 1     | 0.285505 | 1.001694    | -1.426171           | 0.0      |
| 2     | 0.285505 | 0.161103    | 0.489988            | 0.0      |
| 3     | 0.245062 | -0.492690   | 0.416289            | 0.0      |
| 4     | 0.210471 | -0.679488   | 0.416289            | 0.0      |
| ...   | ...      | ...         | ...                 | ...      |
| 99995 | 0.285505 | 0.628098    | -1.180510           | 0.0      |
| 99996 | 0.181524 | 0.908295    | -0.934848           | 0.0      |
| 99997 | 0.290835 | 0.161103    | 0.416289            | 0.0      |
| 99998 | 0.370154 | -1.426680   | -0.934848           | 0.0      |
| 99999 | 0.234403 | 1.001694    | -1.180510           | 0.0      |

[99996 rows x 9 columns]

## Extract The Features Using SelectFromModel imported from sklearn package

```
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
```

```

from sklearn.feature_selection import SelectFromModel

X = df.drop(columns=['diabetes'])
y = df['diabetes']

categorical_features = ['gender', 'smoking_history']

numerical_features = ['age', 'hypertension', 'bmi', 'HbA1c_level',
'blood_glucose_level', 'heart_disease']

categorical_preprocessor = OneHotEncoder(handle_unknown='ignore')

numerical_preprocessor = SimpleImputer(strategy='median')

preprocessor = ColumnTransformer(
 transformers=[
 ('num', numerical_preprocessor, numerical_features),
 ('cat', categorical_preprocessor, categorical_features)
])

rf_clf = RandomForestClassifier()
selector = SelectFromModel(estimator=rf_clf)

pipeline = Pipeline([
 ('preprocessor', preprocessor),
 ('selector', selector)
])

pipeline.fit(X, y)
selected_features_indices =
pipeline.named_steps['selector'].get_support(indices=True)

selected_features = df.columns[selected_features_indices]

print("The Selected Featurs are : ",selected_features)

```

```

The Selected Features are: Index(['gender', 'hypertension', 'heart_disease', 'smoking_history'], dtype='object')

```

Divide the Train and Test data using the train\_test\_split method imported from sklearn

```

from sklearn.model_selection import train_test_split

```

```
X_selected = df[selected_features]
y = df['diabetes']

X_train, X_test, y_train, y_test = train_test_split(X_selected, y,
test_size=0.2, random_state=42)

print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of y_test:", y_test.shape)
```

```
Shape of X_train: (79996, 4)
Shape of X_test: (20000, 4)
Shape of y_train: (79996,)
Shape of y_test: (20000,)
```

#### Learning Outcome:

- 1) Learning the python packages such as Numpy, Scipy, Scikitlearn, Matplotlib
- 2) Learning about Loading the dataset using the pandas
- 3) Learning about the Pre processing steps in ML
- 4) Learning about Feature Selection Techniques.
- 5) Learning about Split the data into training, testing and validation sets