

UCS 2612 Machine Learning Laboratory
Assignment 8

Name : Mega V
Reg No : 3122 21 5001 051

AIM :

To develop a python program to diagnose breast cancer using Ensemble Models and to Visualize the features from the dataset and interpret the results obtained by the model using Matplotlib library.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import cross_val_score
from sklearn.metrics import roc_curve, auc
```

□ LOADING THE DATASET :

```
df = pd.read_csv("data.csv")
df.head()
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...	t
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	...	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	...	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	...	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	...	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	...	

5 rows x 33 columns

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                569 non-null    float64
12  radius_se                             569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                               569 non-null    float64
16  smoothness_se                         569 non-null    float64
17  compactness_se                        569 non-null    float64
18  concavity_se                          569 non-null    float64
19  concave points_se                     569 non-null    float64
20  symmetry_se                           569 non-null    float64
21  fractal_dimension_se                  569 non-null    float64
22  radius_worst                          569 non-null    float64
23  texture_worst                         569 non-null    float64
24  perimeter_worst                       569 non-null    float64
25  area_worst                            569 non-null    float64
26  smoothness_worst                     569 non-null    float64
27  compactness_worst                     569 non-null    float64
28  concavity_worst                       569 non-null    float64
29  concave points_worst                  569 non-null    float64
30  symmetry_worst                        569 non-null    float64
31  fractal_dimension_worst               569 non-null    float64
32  Unnamed: 32                           0 non-null      float64
```

dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB

df.describe()

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry
count	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.
mean	3.037183e+07	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799	0.048919	0.
std	1.250206e+08	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720	0.038803	0.
min	8.670000e+03	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000	0.000000	0.
25%	8.692180e+05	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560	0.020310	0.
50%	9.060240e+05	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.061540	0.033500	0.
75%	8.813129e+06	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130700	0.074000	0.
max	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800	0.201200	0.

8 rows x 32 columns

□ PREPROCESSING :

Null Values :

print(df.isnull().sum())

id	0
diagnosis	0
radius_mean	0
texture_mean	0
perimeter_mean	0
area_mean	0
smoothness_mean	0
compactness_mean	0
concavity_mean	0
concave points_mean	0
symmetry_mean	0
fractal_dimension_mean	0
radius_se	0
texture_se	0
perimeter_se	0
area_se	0
smoothness_se	0
compactness_se	0
concavity_se	0
concave points_se	0
symmetry_se	0
fractal_dimension_se	0
radius_worst	0
texture_worst	0
perimeter_worst	0
area_worst	0
smoothness_worst	0
compactness_worst	0
concavity_worst	0
concave points_worst	0
symmetry_worst	0
fractal_dimension_worst	0
Unnamed: 32	569
dtype:	int64

Removing Unamed: 32 column as its completely NaN :

df.drop(df.columns[-1], axis=1, inplace=True)
df.head()

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...	r
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	...	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	...	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	...	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	...	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	...	

5 rows x 32 columns

Converting categorical to numeric values :

df["diagnosis"] = df["diagnosis"].map({"M" : 0, "B": 1})

```
# Standardisation :

numeric_cols = df.columns[2:]

standardiser = StandardScaler()
df[numeric_cols] = standardiser.fit_transform(df[numeric_cols])
df.head()
```


	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...	r
0	842302	0	1.097064	-2.073335	1.269934	0.984375	1.568466	3.283515	2.652874	2.532475	...	
1	842517	0	1.829821	-0.353632	1.685955	1.908708	-0.826962	-0.487072	-0.023846	0.548144	...	
2	84300903	0	1.579888	0.456187	1.566503	1.558884	0.942210	1.052926	1.363478	2.037231	...	
3	84348301	0	-0.768909	0.253732	-0.592687	-0.764464	3.283553	3.402909	1.915897	1.451707	...	
4	84358402	0	1.750297	-1.151816	1.776573	1.826229	0.280372	0.539340	1.371011	1.428493	...	

5 rows x 32 columns

□ EXPLORATORY DATA ANALYSIS :


```
# Heat Map :

corr_matrix = df.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, cmap='viridis', center=0, square=True, linewidths=.5)
plt.title('Correlation Heatmap')
plt.show()
```



```
# Box Plot :
```

```
plt.figure(figsize=(10, 8))
sns.boxplot(data=df[["radius_mean", "area_mean", "symmetry_mean", "area_worst", "fractal_dimension_worst"]])
plt.title('Box Plot of Features')
plt.xlabel('Features')
plt.show()
```



```
print(df.shape)
```

```
(569, 32)
```

```
# Outlier Detection :
```

```
z_scores = stats.zscore(df)
threshold = 4
outlier_mask = (z_scores > threshold) | (z_scores < -threshold)
df = df[~outlier_mask.any(axis=1)]
```

```
print(df.shape)
```

```
(526, 32)
```

□ TRAINING AND TESTING SPLIT :

```
X = df.drop("diagnosis", axis = 1)
Y = df["diagnosis"]

X_train, X_test, y_train, y_test = train_test_split(X, Y, random_state = 42)
```

□ RANDOM FOREST MODEL :

```
rf_classifier = RandomForestClassifier(n_estimators=20, random_state=42)
rf_classifier.fit(X_train, y_train)
y_pred = rf_classifier.predict(X_test)

train_accuracy = accuracy_score(y_train, rf_classifier.predict(X_train))
test_accuracy = accuracy_score(y_test, y_pred)

print("***** Random Forest *****")
print("\nTraining Accuracy:", round(train_accuracy, 4))
print("Test Accuracy:", round(test_accuracy, 4))
```

```
***** Random Forest *****
```

```
Training Accuracy: 0.9924
```

Test Accuracy: 0.9697

```
# Cross Validation :
```


```
cv_scores = cross_val_score(rf_classifier, X, Y, cv=5)
print("Cross-Validation Score:", round(cv_scores.mean(), 4))
```

Cross-Validation Score: 0.9468

```
# ROC Curve :
```

```
y_pred_proba = rf_classifier.predict_proba(X_test)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



As the cross validation score is 94%, the Random Forest ensemble model is not overfitting and hence performing better with an training accuracy of 99% and testing accuracy of 97%

□ ADABOOST ENSEMBLE MODEL :

```
adaboost_classifier = AdaBoostClassifier(n_estimators=20, random_state=42)
adaboost_classifier.fit(X_train, y_train)
y_pred = adaboost_classifier.predict(X_test)

train_accuracy = adaboost_classifier.score(X_train, y_train)
test_accuracy = accuracy_score(y_test, y_pred)

print("***** ADABOOST *****")
print("\nTraining Accuracy:", round(train_accuracy, 4))
print("Testing Accuracy:", round(test_accuracy, 4))
```

***** ADABOOST *****

Training Accuracy: 1.0
Testing Accuracy: 0.9697

c:\Python312\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will
warnings.warn(
c:\Python312\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will
warnings.warn(
c:\Python312\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will

```
# Cross Validation :
```

```
cv_scores = cross_val_score(adaboost_classifier, X, Y, cv=5)
print("Cross-Validation Score:", round(cv_scores.mean(), 4))
```


c:\Python312\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will
warnings.warn(
c:\Python312\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will
warnings.warn(
c:\Python312\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will

```
warnings.warn(
c:\Python312\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will
warnings.warn(
c:\Python312\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:519: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will
warnings.warn(
Cross-Validation Score: 0.9639
```

ROC Curve :

```
y_pred_proba = adaboost_classifier.predict_proba(X_test)[: , 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



As the cross validation score is 96%, the Adaboost ensemble model is not overfitting and hence performing better with an training accuracy of 100% and testing accuracy of 96%

...					
		Training Accuracy	Testing Accuracy	Cross Validation Score	AUC
Random Forest Model :		99%	97%	95%	99%
Adaboost :		100%	96%	96%	100%
...					
'\n		Training Accuracy	Testing Accuracy	Cross Validation Score	AUC\n\nRandom
Forest Model :		99%	97%	95%	99%\n\nAdaboost
:	100%	96%	96%	100%\n\n'	

Both models are performing good without overfitting, But Adaboost is slightly better in training whereas Random forest is slightly better in testing phase. Overall, Adaboost has a slight edge over Random Forest in terms of training accuracy and cross validation score making it perform better for unseen data.

Learning Outcome :

1. Interpret the results and evaluate the performance of random forest models.
2. Implement and tune random forest algorithms for predictive modeling tasks. Apply executable techniques to improve the performance of machine learning models.